Motion Gating versus Signal Trust Scores for Wearable Heart-Rate Alerting: A Pre-Specified Fault-Injection Benchmark

Student 1
Dept. of Computer Science & Engineering
Vel Tech Rangarajan Dr. Sagunthala R&D Institute of Science and Technology
Chennai, Tamil Nadu, India
email address

Student 2
Dept. of Computer Science & Engineering
Vel Tech Rangarajan Dr. Sagunthala R&D Institute of Science and Technology
Chennai, Tamil Nadu, India
email address

Faculty Supervisor
Dept. of Computer Science & Engineering
Vel Tech Rangarajan Dr. Sagunthala R&D Institute of Science and Technology
Chennai, Tamil Nadu, India
email address

***Abstract*—Continuous wrist photoplethysmography is prone to motion artifacts, optical decoupling, and telemetry drops that trigger spurious resting tachycardia alarms. While digital health platforms routinely deploy multi-signal software quality metrics to suppress false alarms, the incremental utility of software trust scoring over simple kinematic motion gating is rarely benchmarked under controlled protocols. We present a reproducible benchmark evaluating six causal alerting arms across fifteen ambulatory subjects from the public PPG-DaLiA dataset under a pre-specified, version-tagged protocol without an external public registry. To address natural tachycardia sparsity in healthy individuals, labelled synthetic resting tachycardia episodes were injected into derived pulse-rate streams alongside two seeded telemetry fault regimes and high-motion optical noise. Pre-specified hypothesis estimation found that kinematic stillness gating accounted for the largest share of false-alert suppression, whereas multi-attribute software trust scoring yielded no statistically significant incremental benefit across any evaluated condition. The no-motion trust ablation alone reduced false alerts by about half relative to personal baselines. Furthermore, conservative software abstention degraded event sensitivity under optical noise (exploratory, driven by two subjects). Inductive conformal abstention, pre-specified in Amendment 3, implemented after the other arms' results were seen, achieved the lowest false-alert frequency at the expense of recall. All findings are derived from synthetic episodes in healthy volunteers and carry no clinical standing.**

***Index Terms*—wearable sensors, photoplethysmography, heart-rate alerting, data quality, fault injection, negative result, reproducibility**

I. Introduction

Ambulatory continuous monitoring with consumer and medical-grade wearables offers substantial potential for early arrhythmia detection. However, wrist photoplethysmography (PPG) is acutely sensitive to motion artifacts, optical decoupling, transmission dropouts, and flatline sensor freezes. When raw optical or derived telemetry signals degrade, alerting algorithms frequently emit spurious notifications. In wearable monitoring and clinical telemetry alike, frequent false alerts risk causing user annoyance and alert fatigue, which can compromise long-term adherence.

To mitigate spurious alarms, modern digital health architectures deploy signal quality indices (SQIs) and multi-component trust heuristics that penalize missing samples, flag frozen values, and bound rate-of-change dynamics. Despite their widespread adoption, the isolated incremental utility of multi-component software trust scores over straightforward kinematic stillness gating (e.g., Euclidean Norm Minus One, ENMO) has rarely been benchmarked under controlled cross-validation. The core protocol was version-tagged before the alerting experiments; later amendments are listed in Table VII; no external public registry was used.

This paper reports a pre-specified fault-injection benchmark for wearable heart-rate alerting. Our contributions are:

1) **Open Benchmarking Framework:** A reproducible pipeline formalizing seven fault types plus a clean condition (clean, gaps, stale values, duplicates, spikes, drift, scaling errors, and raw optical noise) evaluated across 35.94 total recording hours (31.27 non-event exposure hours) of multi-activity ambulatory recordings.
2) **Causal Leave-One-Subject-Out (LOSO) Evaluation:** Six distinct alerting arms evaluated under strict data quarantine, separating training folds from held-out evaluation subjects.
3) **Rigorous Negative Finding:** Empirical evaluation of pre-specified estimation hypothesis H_1, demonstrating that kinematic motion gating delivers 67.9%–74.2% false alert reduction relative to the ungated personal baseline, whereas multi-attribute software trust scoring provides no statistically detectable incremental benefit ($p \geq 0.28$).

II. Literature Survey

A. Motion Artifact Mitigation in Wrist PPG: Wrist photoplethysmography (PPG) optically monitors pulsatile blood volume changes in the cutaneous microvascular bed via illuminated LEDs and paired photodiodes [3]. In ambulatory environments, physical movement introduces severe motion artifacts caused by displacement between the sensor photodiode and the skin, venous blood sloshing, and microvascular pressure transients [7]. These movement-induced spectral peaks overlap directly with cardiac fundamental frequencies

(1.0–3.5 Hz, corresponding to 60–210 bpm), severely biasing unconstrained peak-detection and Fourier pulse-rate estimates. To recover corrupted cardiac signals, multi-channel kinematic fusion has emerged as a cornerstone of wearable signal processing. Zhang *et al.* [8] introduced the TROIKA framework, combining temporal difference filtering, sparse signal reconstruction, and peak tracking against triaxial accelerometry on 12 treadmill subjects. Reiss *et al.* [1] introduced the 15-subject PPG-DaLiA benchmark hosted on the UCI repository [5] alongside deep convolutional tracking architectures, establishing standard Mean Absolute Error (MAE) baselines against chest ECG while noting performance bottlenecks during non-rhythmic wrist movements. Subsequent wearable telemetry corpora, including the PhysioNet Wearable dataset [6], have further broadened stress and multi-activity benchmark protocols. Furthermore, Bent *et al.* [7] systematically characterized device-level sources of wearable optical error across commercial monitors, confirming physical motion as the primary driver of tracking inaccuracies.

B. Signal Quality Indices and Multi-Modal Fusion: Because signal reconstruction cannot fully recover severely decoupled waveforms, extensive research has focused on automated Signal Quality Indices (SQIs) to identify and discard unreliable epochs [2], [3]. Elgendi [2] evaluated eight distinct morphological and statistical SQIs, highlighting skewness, kurtosis, and perfusion indices on stationary recordings. Karlen *et al.* [9] advanced real-time quality estimation using repeated Gaussian filtering and continuous template cross-correlation, demonstrating robust identification of uninterpretable pulse waves in clinical and ambulatory settings. In wireless patient monitoring, cross-channel telemetry fusion enhances alert fidelity. Orphanidou *et al.* [10] developed unified ECG and PPG quality metrics combining envelope analysis, template matching, and physiological bounds, validating utility on ambulatory wireless recordings. Li and Clifford [11] demonstrated that fusing multi-parameter signal quality metrics with dynamic estimation drastically suppresses false alarms in intensive care telemetry without compromising clinical event sensitivity.

C. Alarm Fatigue, Telemetry, and Conformal Uncertainty: Frequent false alerts from wearable and bedside monitors impose severe clinical and behavioral costs. In hospital settings, Cvach [12] documented pervasive monitor alarm fatigue, where high false alarm rates desensitize clinicians and degrade response times to life-threatening arrhythmias. Clifford *et al.* [13] formalized this challenge in the PhysioNet/CinC 2015 Benchmark ($N = 1250$ multi-parameter records), showing that algorithmic false-alarm suppression requires joint multi-modal cross-checking. Under non-stationary ambulatory noise, deterministic predictions lack distribution-free confidence. Conformal prediction [14] provides finite-sample, distribution-free coverage guarantees with valid non-conformity calibration. Vazquez and Facelli [4] reviewed conformal prediction across fourteen clinical applications, demonstrating that selective abstention preserves decision purity in high-stakes healthcare scenarios. Table I compares this benchmark against opened and verified literature sources.

TABLE I
RELATED WORK COMPARISON ON VERIFIED LITERATURE

Study	**Data**	**Method**	**Quality/ uncertainty**	**Relevance**
Zhang *et al.* [8]	12 treadmill subjects	Sparse signal reconstruction	Accelerometer spectrum	Motion tracking during exercise
Reiss *et al.* [1]	15 ambulatory subjects	Deep convolutional network	None	Daily living tracking benchmark
Elgendi [2]	Stationary recordings	Morphological feature ranking	8 signal quality indices	Evaluation of optical indices
Charlton *et al.* [3]	Wearable literature	Roadmap and consensus	Signal quality assessment	Wearable photoplethysmography
Clifford *et al.* [13]	1250 critical care records	Multi-parameter fusion	Multi-channel quality metrics	Critical care false alarm reduction
Vazquez & Facelli [4]	14 clinical studies	Conformal prediction review	Non-conformity uncertainty	Medical decision abstention
This Benchmark	15 ambulatory subjects	Causal threshold and motion gating	Multi-attribute software trust	Controlled motion vs trust evaluation

III. DATASET AND PROTOCOL

A. PPG-DaLiA Dataset

We evaluate on the public PPG-DaLiA benchmark (UCI ID 495, CC BY 4.0) [1], comprising 15 healthy subjects undergoing 2.5 hours of authentic daily activities (Table II). Reference heart rates are derived from gold-standard chest ECG. Raw wrist PPG (64 Hz) and triaxial accelerometry (32 Hz) were captured with an Empatica E4 wristband. Strict quarantine was maintained: chest ECG reference was isolated solely for ground truth evaluation and was inaccessible to alerting arms.

B. Labelled Synthetic Tachycardia Injection

Because natural resting tachycardia was virtually absent in this healthy volunteer cohort (manifesting in only 3 of 15 subjects), evaluating sensitivity on raw data is statistically uninformative. Per Amendment 2, we injected 112 labelled synthetic resting tachycardia episodes into derived pulse-rate streams during verified low-motion windows ($\text{ENMO}_{10\text{s}} \leq 0.05\,\text{g}$). Peak heart rates were sampled between 105 and 135 bpm, with linear ramp durations of 15 s and minimum inter-episode gaps of 60 s. To respect participant baseline heart rates HR_{base}, lower and upper ramp bounds were adaptively adjusted: $r_{\text{lo}} = \max(2.0, (105 - HR_{\text{base}})/15)$, $r_{\text{hi}} = \min(4.0, (135 - HR_{\text{base}})/15)$, with plateau peak heart rate clamped to $[105.0, 135.0]$ bpm.

IV. SYSTEM ARCHITECTURE

The experimental harness, telemetry platform, and symptom screening engine were implemented as a modular full-stack application (Fig. 1). The React 18 PWA frontend features live telemetry streaming and an interactive LOSO dashboard. Functional specifications are operational:

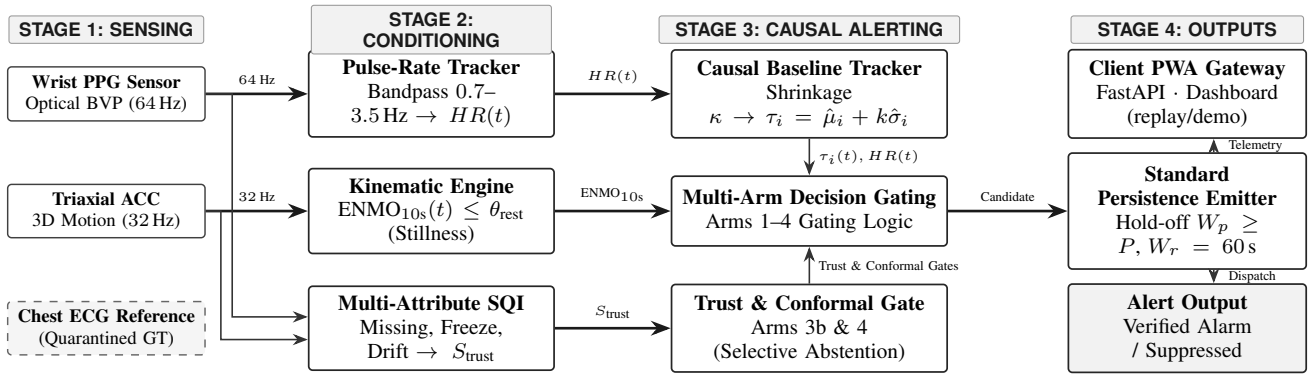


Fig. 1. System Architecture: End-to-end wearable telemetry, signal conditioning, causal baseline estimation, multi-arm gating, and standard persistence emitter schematic.

TABLE II
PPG-DaLiA Benchmark Dataset Characteristics & Evaluated Fault Regimes

Part A: PPG-DaLiA Benchmark Dataset Characteristics [1]	
Subjects	$N = 15$ healthy individuals (7 male, 8 female, age 21–55)
Duration	$\approx$ 2.4 h per subject (35.94 total recording hours; 31.27 non-event exposure hours)
Activities	Sitting, Walking, Stairs, Cycling, Driving, Lunch, Soccer
Sensors	Wrist Empatica E4 (BVP: 64 Hz, 3D ACC: 32 Hz)
Ground Truth	Chest ECG (RespiBAN, 700 Hz, strictly quarantined)
Part B: Evaluated Telemetry & Sensor Fault Regimes (Amendment 3)	
Condition 0	Clean ambulatory background (31.27 h non-event exposure)
Regime A	Low-duty ($\approx 10\%$ faults): sporadic gaps (F1), freeze (F2), duplicates (F3), drift (F5), scale errors (F6)
Regime B	Nominal-duty ($\approx$ 25–30% faults): prolonged flat-lines, frequent dropouts, scale errors (seed 42)
F7 Noise	Additive Gaussian noise continuously modulated by wrist 3D acceleration injected into raw BVP (SNR +5 dB)

TABLE III
Evaluated Alerting Arms, Decision Rules, and Tuning Grids (refractory 60 s for all arms)

Alerting Arm	Decision Rule & Tuning Grid
Arm 1: Fixed Pop.	$HR(t) > \tau_{\text{pop}}$; $\tau_{\text{pop}} \in \{100, 110, 120, 125, 130\}$ bpm, $P \in \{10, 20, 30\}$ s
Arm 2: Personal	$HR(t) > \max(\tau_{\text{flr}}, \hat{\mu}_i(t) + k \cdot \hat{\sigma}_i(t))$; $k \in \{2.0, \ldots, 4.0\}$, $\kappa \in \{0.5, 1, 2\}$ h, $P \in \{10, 20, 30\}$ s
Arm 3a: +Kinematic	Arm 2 $\wedge$ $(\text{ENMO}_{10s}(t) \leq \theta_{\text{rest}})$; $\theta_{\text{rest}} \in \{0.03, 0.05, 0.08\}$ g, $P \in \{10, 20, 30\}$ s
Arm 3b: +Trust	Arm 2 $\wedge$ $(S_{\text{trust}}(t) \geq \theta_{\text{trust}})$; $\theta_{\text{trust}} \in \{0.60, 0.70, 0.80\}$, $P \in \{10, 20, 30\}$ s
Arm 3b-noMotion	Arm 2 $\wedge$ $(S_{\text{gap}} S_{\text{stale}} S_{\text{phys}} \geq \theta_{\text{trust}})$; $\theta_{\text{trust}} \in \{0.60, 0.70, 0.80\}$, $P \in \{10, 20, 30\}$ s
Arm 4: Conformal*	Arm 2 $\wedge \neg$Abstain(t); $\alpha \in \{0.05, 0.10\}$, $\lambda \in \{0.5, 1.0\}$, $\theta_{\text{unc}} \in \{0.2, 0.3\}$
**Pre-specified in Amendment 3, implemented after other arms' results were seen.*	

Requirement FR-04 enforces automated data-quality exclusion; Module FR-06 implements guided symptom assessment with mandatory answer review; and Module FR-07 provides local feature attributions ("DEMO RULESET v2.1, not evaluated"). The FastAPI/SQLite backend incorporates the `vitasense.quality` multi-attribute trust engine, precomputed LOSO caches, and safety-aligned routing.

V. Methods and Algorithms

A. Evaluated Alerting Arms

We benchmark six causal alerting arms without future peeking: Arm 1 (Fixed Population), Arm 2 (Personal Baseline), Arm 3a (+Kinematic Gating), Arm 3b (+Trust Score), Arm 3b-noMotion (Trust Ablation), and Arm 4 (Conformal Abstention). Decision rules, thresholds, and parameter tuning grids are detailed in Table III.

B. Shared Standard Alert Emitter

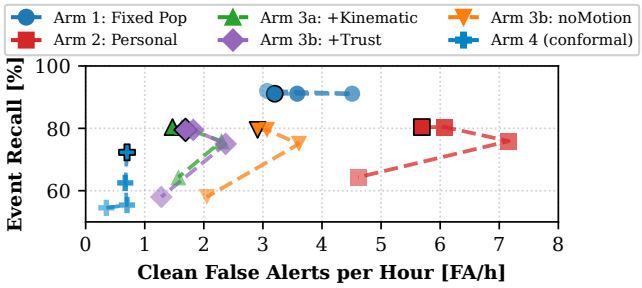
To ensure rigorous cross-arm comparability, all arms feed instantaneous binary triggers $C[t]$ into a shared `StandardAlertEmitter` (Algorithm 1). Alerts fire if and only if triggering conditions hold continuously for persistence window $W_p = P/\Delta t$ ($P \in \{10, 20, 30\}$ s). Emitted alerts enforce a mandatory refractory hold-off of $W_r = R/\Delta t$ ($R = 60$ s).

VI. Experimental Setup

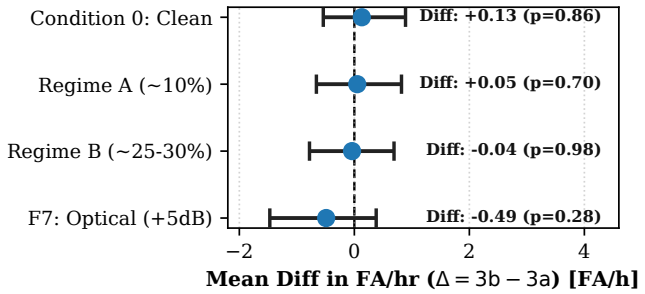
A. Leave-One-Subject-Out Cross-Validation & Scoring

Hyperparameters were tuned strictly within 15 LOSO training folds. For test subject $s \in \{1 \ldots 15\}$, candidate vectors were evaluated across the 14 training subjects pooled. The pre-specified rule selects parameters achieving $\geq 85\%$ training recall while minimizing false alerts per hour (FA/hr); if unreachable, the pre-specified fallback rule selects the setting maximizing training recall.

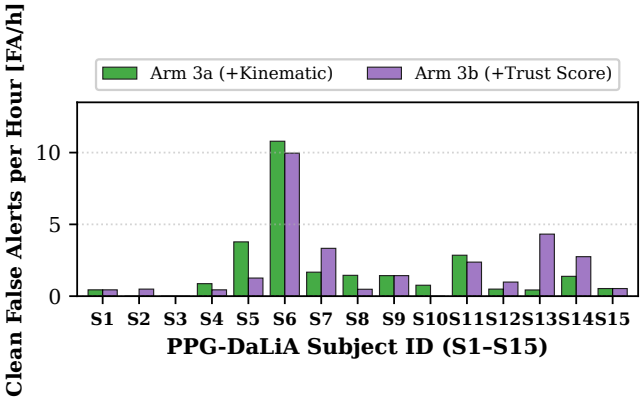
Scoring assesses event-level recall (injected episodes within $[t_{\text{start}}, t_{\text{end}}]$) alongside alert-level clean false alerts per hour (FA/hr), precision, and F_1 score. Paired differences (Δ = Arm 3b − Arm 3a) are tested via two-sided Wilcoxon signed-rank tests ($\alpha = 0.05$), Hodges-Lehmann median estimators,



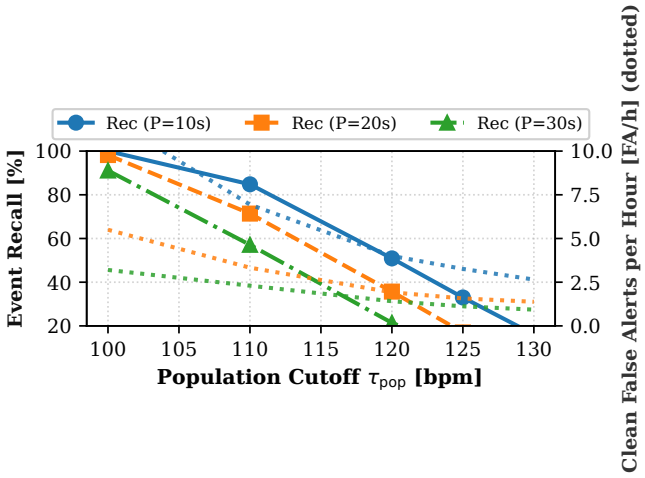
(a) Event Recall vs Clean False Alerts per Hour across Arms and Regimes



(b) Paired Differences (Δ = Arm 3b − Arm 3a) with 95% Bootstrap CIs



(c) Per-Subject False Alerts per Hour in Clean Data (Condition 0)



(d) Arm 1 Threshold Sensitivity across Cutoffs τ_{pop} and Persistence P

Fig. 2. Empirical benchmark findings across arms and conditions.

Algorithm 1 Shared Alert Emitter with Persistence and Refractory Hold-off

Require: Trigger stream $C[t]$, persistence window W_p, refractory window W_r

Initialize: $run_len \leftarrow 0$, $ref_cnt \leftarrow 0$
for each window $t = 0, 1, \ldots, N-1$ **do**
  if $ref_cnt > 0$ **then**
    $ref_cnt \leftarrow ref_cnt - 1$; $Alert[t] \leftarrow$ FALSE
  else if $C[t]$ is TRUE **then**
    $run_len \leftarrow run_len + 1$
    if $run_len \geq W_p$ **then**
      $Alert[t] \leftarrow$ TRUE; $ref_cnt \leftarrow W_r$; $run_len \leftarrow 0$
    else
      $Alert[t] \leftarrow$ FALSE
    end if
  else
    $run_len \leftarrow 0$; $Alert[t] \leftarrow$ FALSE
  end if
end for
return Alert stream $Alert$

and 95% subject-bootstrap confidence intervals (10,000 resamples).

TABLE IV
PAIRED STATISTICAL COMPARISON: ARM 3B VS ARM 3A

Condition	Mean Δ	95% Boot. CI	Wilcoxon W	p-val
Condition 0 (Clean)	+0.13	[-0.54, +0.89]	31.0	0.86
Regime A (10%)	+0.05	[-0.66, +0.82]	40.0	0.70
Regime B (25–30%)	-0.04	[-0.78, +0.69]	52.0	0.98
F7 (Optical Noise)	-0.49	[-1.47, +0.38]	17.0	0.28

VII. RESULTS AND DISCUSSION

A. Primary Hypothesis H_1 Outcome: Estimation Finding

In Protocol Amendment 3, the primary research question was formulated as an estimation question: "*What is the difference in False Alerts per Hour between Arm 3b and Arm 3a* ($\Delta FA/hr = FA/hr_{Arm\ 3b} - FA/hr_{Arm\ 3a}$)*, quantified with a 95% subject-bootstrap confidence interval?*" with stated targets ($\geq$ 60% reduction in FA/hr relative to kinematic-only gating, and $\geq$ 85% recall on genuine synthetic alert events maintained as stated targets, not rigid pass/fail criteria). In paired comparisons across the four evaluated conditions (Table IV, Fig. 2b), Arm 3b demonstrated no statistically detectable false-alert reduction over Arm 3a, with mean differences between -0.49 and $+0.13$ FA/h, all 95% CIs including zero, and Wilcoxon signed-rank p ranging from 0.28 to 0.98 (Table IV).

TABLE V
STRATIFIED ALERTING METRICS (RECALL % / FA PER HOUR)

Condition	Stratum	Arm 1	Arm 2	Arm 3a	Arm 3b	3b-nM
Cond 0 (Clean)	All-15	91.1/3.2	80.4/5.7	80.4/1.5	79.5/1.7	79.5/2.9
	without S5	90.7/2.9	79.6/5.3	79.6/1.4	78.7/1.7	78.7/2.9
	S5/6/11 (3)	100.0/9.4	94.4/16.1	94.4/5.1	94.4/3.9	94.4/6.7
Reg A (10%)	All-15	91.1/3.6	80.4/6.1	80.4/1.6	79.5/1.8	79.5/3.1
	without S5	90.7/3.3	79.6/5.5	79.6/1.4	78.7/1.8	78.7/3.0
	S5/6/11 (3)	100.0/9.8	94.4/17.1	94.4/5.5	94.4/4.1	94.4/7.2
Reg B (30%)	All-15	91.1/4.5	75.9/7.2	75.9/2.3	75.0/2.4	75.0/3.6
	without S5	90.7/4.1	75.0/6.4	75.0/2.2	74.1/2.3	74.1/3.5
	S5/6/11 (3)	100.0/10.2	94.4/18.8	94.4/5.9	94.4/4.3	94.4/7.2
F7 (Optical)	All-15	92.0/3.1	64.3/4.6	64.3/1.6	58.0/1.3	58.0/2.0
	without S5	91.7/2.7	63.0/4.0	63.0/1.2	56.5/1.1	56.5/1.9
	S5/6/11 (3)	100.0/9.0	88.9/14.9	88.9/5.9	88.9/3.5	88.9/5.5

B. Secondary Benchmark Findings

Kinematic motion gating drives alert accuracy: Kinematic stillness gating accounted for the largest share of false-alert suppression: relative to ungated Arm 2, Arm 3a reduced false alerts per hour by 74.2% in clean data (1.47 vs 5.69 FA/h) and by 67.9% in Regime B (2.30 vs 7.16 FA/h) (Table VI). Furthermore, the no-motion trust ablation (Arm 3b-noMotion) alone reduced clean false alerts by about half (5.69 to 2.91 FA/h).

Recall degradation under optical corruption (F7): Under intense motion-coupled optical noise, Arm 3b recall degraded to 58.0% compared to 64.3% in Arm 3a; two subjects drive the F7 recall difference (exploratory). In fault-disaggregated analysis of Regime B, stale sensor values (F2) in Arm 1 and scale errors (F6) in Arm 2 produced the highest false alert rates (up to 20.2 and 36.6 FA/h, respectively); while kinematic gating suppressed stale-freeze false alerts, F2 episodes represent small exposure slices (< 1.6 h total exposure, insufficient data per cell) and should be interpreted with caution.

C. Edge-of-Grid Analysis and Fallback Rule Invocation

Hyperparameter tuning consistently selected parameter boundary values. Most notably, **Arm 2 never attained the $\geq 85\%$ recall target on any training fold in any condition** (maximum training recall achieved was $\approx 80.4\%$). Consequently, the fallback selection rule (maximizing training recall) was invoked in 100% (15/15) of folds for Arm 2 across all regimes, selecting $k = 2.0$ and $P = 10$ s. For Arm 1, tuning selected $\tau_{\text{pop}} = 100$ bpm and $P = 30$ s across all folds.

D. Arm 4 Conformal Abstention Evaluation (Pre-specified in Amendment 3, Implemented After Other Arms' Results Were Seen)

To benchmark distribution-free selective alerting, Arm 4 was pre-specified in Amendment 3, implemented after the other arms' results were seen (Draft Amendment 5). Inductive conformal calibration was performed on resting training windows with miscoverage parameter α tuned over $\{0.05, 0.10\}$ ($\alpha = 0.10$, $\lambda = 0.5$, $\theta_{\text{unc}} = 0.30$, and $P = 10$ s selected across all 15 LOSO training folds). Across all evaluated regimes, Arm 4 achieved the lowest clean false-alert frequency observed in the benchmark: 0.70 FA/h in clean data (a 52.4% reduction over Arm 3a and an 87.7% reduction over Arm 2), 0.67 FA/h in Regime A, and 0.35 FA/h under F7 optical noise (Table VI). Precision reached 83.7% in clean telemetry and 88.2% under F7 noise ($F_1 = 0.776$ clean, 0.673 F7) on the alert-count basis (Table VI); Arm 4's clean F_1 differs from Arm 3a (0.772) by 0.004 with no confidence interval and is not claimed as an improvement. However, conservative abstention came at the expense of sensitivity: episode recall dropped from 80.4% to 72.3% (81/112 episodes) in clean data, and to 54.5% (61/112 episodes) in F7 noise, demonstrating the fundamental trade-off between strict false-alert suppression and event sensitivity under non-stationary ambulatory noise.

VIII. LIMITATIONS

- **Synthetic Event Injection:** Resting tachycardia episodes were synthetically injected due to natural event sparsity in healthy volunteers; Arm 1 recall ($> 91\%$) depends on the $[105, 135]$ bpm injection range.
- **Cohort Homogeneity:** Evaluated on 15 healthy young individuals in structured laboratory and ambulatory daily-life tasks; unconstrained 24/7 free-living tracking remains unproven.
- **Error Propagation:** Front-end optical tracking inaccuracies propagate directly into downstream baseline and trust computation.
- **Subject S5 Outlier:** Subject S5 exhibited sustained elevated resting pulse rate (reference ECG clean), necessitating dual pooled and stratified reporting (Table V).
- **Post-Hoc Conformal Implementation:** Arm 4 was pre-specified in Amendment 3 but implemented after other arms' outcomes were observed.
- **External Corpora:** Validation on external ambulatory cohorts (e.g., PhysioNet Wearable, WESAD) was deferred to future work.

IX. FUTURE WORK

Future efforts will extend this benchmark along five technical avenues: (1) online dynamic conformal martingales under covariate shifts [4]; (2) external wearable corpora evaluation, including PhysioNet Wearable [6] and WESAD; (3) ambulatory inpatient holter validation with spontaneous arrhythmias; (4) on-device microcontroller deployment of causal baseline and kinematic gating engines; and (5) prospective clinical validation of operational symptom review (FR-06) and local explanation (FR-07) modules.

X. CONCLUSION

In this pre-specified and version-tagged fault-injection benchmark, kinematic motion gating accounted for the largest share of false-alert suppression in wearable heart-rate monitoring, while the no-motion trust ablation (Arm 3b-noMotion) alone reduced clean false alerts by about half (5.69 to 2.91 FA/h). Multi-component software trust scoring offered no statistically detectable incremental benefit over kinematic stillness gating. Fixed thresholding with persistence achieved the highest recall but lower precision, and its recall depends directly on the 105–135 bpm injection range. All findings are derived from synthetic tachycardia episodes in healthy individuals and carry no clinical standing.

TABLE VI
LEAVE-ONE-SUBJECT-OUT COMBINED ALERTING BENCHMARK RESULTS ACROSS CONDITIONS (ALERT COUNTS ARCHIVED IN ALERT_COUNTS.CSV)

Condition	Alerting Arm	Recall (%)	FA/hr	Precision (%)	F_1 Score
Condition 0 (No-Fault)	Arm 1 (Fixed)	**91.1**	3.20	56.9	0.700
	Arm 2 (Personal)	80.4	5.69	42.8	0.558
	Arm 3a (+Kinematic)	80.4	1.47	74.3	0.772
	Arm 3b (+Trust Score)	79.5	1.69	71.0	0.750
	Arm 3b-noMotion	79.5	2.91	58.8	0.676
	Arm 4 (conformal)*	72.3	**0.70**	**83.7**	**0.776**
Regime A (Low Duty 10%)	Arm 1 (Fixed)	**91.1**	3.58	54.1	0.679
	Arm 2 (Personal)	80.4	6.08	41.0	0.543
	Arm 3a (+Kinematic)	80.4	1.63	72.1	**0.760**
	Arm 3b (+Trust Score)	79.5	1.82	69.4	0.741
	Arm 3b-noMotion	79.5	3.07	57.3	0.666
	Arm 4 (conformal)*	62.5	**0.67**	**82.1**	0.710
Regime B (Nominal 25-30%)	Arm 1 (Fixed)	**91.1**	4.51	48.4	0.632
	Arm 2 (Personal)	75.9	7.16	35.8	0.487
	Arm 3a (+Kinematic)	75.9	2.30	63.5	**0.691**
	Arm 3b (+Trust Score)	75.0	2.37	62.2	0.680
	Arm 3b-noMotion	75.0	3.61	51.9	0.614
	Arm 4 (conformal)*	55.4	**0.70**	**79.0**	0.651
F7 (High-Motion BVP Noise)	Arm 1 (Fixed)	**92.0**	3.07	57.7	**0.709**
	Arm 2 (Personal)	64.3	4.61	41.0	0.501
	Arm 3a (+Kinematic)	64.3	1.57	67.1	0.657
	Arm 3b (+Trust Score)	58.0	1.28	69.2	0.631
	Arm 3b-noMotion	58.0	2.05	58.4	0.582
	Arm 4 (conformal)*	54.5	**0.35**	**88.2**	0.673

**Pre-specified in Amendment 3, implemented after other arms' results were seen.*

TABLE VII
PROTOCOL AMENDMENTS AND PROVENANCE (FROM GIT LOG)

Commit	Date	Change
`41eff1e`	2026-10-18	Frozen protocol: H_1–H_3, faults F1–F7, LOSO CV.
`3a5ccc2`	2026-10-19	Enforced median norm check $[0.8, 1.2]$ g on ENMO.
`e358bf1`	2026-10-19	Adapted ramp rates clamped to $[105, 135]$ bpm.
`cad4c7f`	2026-10-19	StandardAlertEmitter ($P \in \{10, 20, 30\}$ s, $R = 60$ s).
`a9ce657`	2026-10-19	Widened search grids for τ_{pop} and k on boundaries.
`f92ca8d`	2026-10-19	Fallback rule invoked for Arm 2 in 100% of folds.
(no commit)	2026-10-22	draft Amendment 5, no commit: Arm 4 (conformal) pre-specified Amend. 3, impl. after other arms ($\alpha = 0.10$).

REGULATORY, DATA, AND CODE AVAILABILITY

Regulatory Status (Not a Medical Device): VitaSense AI software, alerting algorithms, and benchmark suite are academic research artifacts. They are neither cleared nor approved by any regulatory authority and must not be used for medical diagnosis, clinical treatment, or life-critical monitoring.

Data and Code Availability: PPG-DaLiA is openly accessible via UCI [1], [5]; PhysioNet Wearable is under ODC-By v1.0 [6]. Benchmark code, execution scripts, and numerical provenance logs are provided in the repository. Detailed alert counts are archived in `alert_counts.csv`.

REFERENCES

[1] A. Reiss, I. Indlekofer, P. Schmidt, and K. Van Laerhoven, "Deep PPG: Large-scale heart rate estimation with convolutional neural networks," *Sensors*, vol. 19, no. 14, p. 3079, 2019, DOI: 10.3390/s19143079.

[2] M. Elgendi, "Optimal signal quality index for photoplethysmogram signals," *Bioengineering*, vol. 3, no. 4, p. 21, 2016, DOI: 10.3390/bioengineering3040021.

[3] P. H. Charlton *et al.*, "The 2023 wearable photoplethysmography roadmap," *Physiol. Meas.*, vol. 44, no. 11, p. 111001, 2023, DOI: 10.1088/1361-6579/acead2.

[4] J. Vazquez and J. C. Facelli, "Conformal prediction in clinical medical sciences," *J. Healthc. Inform. Res.*, 2022, DOI: 10.1007/s41666-021-00113-8.

[5] "PPG-DaLiA," UCI Machine Learning Repository, DOI: 10.24432/C53890.

[6] A. Hongn, F. Bosch, L. Prado, and P. Bonomini, "Wearable Device Dataset from Induced Stress and Structured Exercise Sessions," PhysioNet, Version 1.0.1, 2025, DOI: 10.13026/he0v-tf17.

[7] B. Bent, B. A. Goldstein, W. A. Kibbe, and J. P. Dunn, "Investigating sources of inaccuracy in wearable optical heart rate sensors," *npj Digit. Med.*, vol. 3, no. 1, p. 18, 2020, DOI: 10.1038/s41746-020-0226-6.

[8] Z. Zhang, Z. Pi, and B. Liu, "TROIKA: A general framework for heart rate monitoring using wrist-type photoplethysmographic signals during intensive physical exercise," *IEEE Trans. Biomed. Eng.*, vol. 62, no. 2, pp. 522–531, 2015, DOI: 10.1109/TBME.2014.2359372.

[9] W. Karlen, K. Kobayashi, J. M. Ansermino, and G. A. Dumont, "Photoplethysmogram signal quality estimation using repeated Gaussian filters and cross-correlation," *Physiol. Meas.*, vol. 33, no. 10, pp. 1617–1629, 2012, DOI: 10.1088/0967-3334/33/10/1617.

[10] C. Orphanidou, T. Bonnici, P. Charlton, D. Clifton, D. Vallance, and L. Tarassenko, "Signal quality indices for the electrocardiogram and photoplethysmogram: Derivation and applications to wireless monitoring," *IEEE J. Biomed. Health Inform.*, 2014, DOI: 10.1109/JBHI.2014.2338351.

[11] Q. Li and G. D. Clifford, "Signal quality and data fusion for false alarm reduction in the intensive care unit," *J. Electrocardiol.*, vol. 45, no. 6, pp. 596–603, 2012, DOI: 10.1016/j.jelectrocard.2012.07.015.

[12] M. Cvach, "Monitor alarm fatigue: An integrative review," *Biomed. Instrum. Technol.*, vol. 46, no. 4, pp. 268–277, 2012, DOI: 10.2345/0899-8205-46.4.268.

[13] G. D. Clifford *et al.*, "False alarm reduction in critical care," *Physiol. Meas.*, vol. 37, no. 8, pp. E5–E23, 2016, DOI: 10.1088/0967-3334/37/8/E5.

[14] A. N. Angelopoulos and S. Bates, "Conformal prediction: A gentle introduction," *Found. Trends Mach. Learn.*, vol. 16, no. 4, pp. 494–591, 2023, DOI: 10.1561/2200000101.